CS107, Lecture 2
Unix, C, Bits and Bytes, Integer Representations

Reading: Bryant & O'Hallaron, Ch. 2.2-2.3 (skim)
Ed Discussion: https://edstem.org/us/courses/70267/discussion/5945594

CS107 Topic 1

How can a computer represent integer numbers?

Why is answering this question important?

• Helps us understand the limitations of computer arithmetic (today and Friday)

• Shows us how to more efficiently perform arithmetic (Friday and Monday)

• Shows us how we can encode data more compactly and efficiently (Monday)

assign1: implement 3 programs that manipulate binary representations to (1) work around the limitations of arithmetic with addition, (2) simulate an evolving colony of cells, and (3) print Unicode text to the terminal.

Learning Goals

- Understand the binary and hexadecimal number systems and how to convert between binary, hexadecimal, and decimal

- Understand how positive and negative numbers are represented in binary

- Learn about overflow, why it occurs, and how overflow can impact program execution

Demo: Unexpected Behavior

```
cp -r /afs/ir/class/cs107/lecture-code/lect02 .
```

Bits

Bits

Computers are built around the idea of two states: "on" and "off". Transistors implement this in hardware, and bits represent this in software.

One Bit At A Time

- We can combine bits, as with base-10 numbers, to represent more data. **8 bits = 1 byte**.

- Computer memory is just a large array of bytes. It is **byte-addressable**, meaning you can't address a bit in isolation, only a full byte.

- Computers still fundamentally operate in bits. It's just that we've gotten more creative about how to represent data.
  - images
  - audio
  - video
  - text

Base 10

5 9 3 4

digits 0 – 9

(or rather, 0 through base – 1)

Base 10

5 9 3 4

↑ thousands ↑ hundreds ↑ tens ↑ ones

= **5** * 1000 + **9** * 100 + **3** * 10 + **4** * 1

Base 10

5 9 3 4

10^3 10^2 10^1 10^0

5 9 3 4

10^X: 3   2   1   0

Base 2

1 0 1 1

2^X: 3   2   1   0

digits 0 – 1

(or rather, 0 through base – 1)

Base 2

$$1 \quad 0 \quad 1 \quad 1$$

$$2^3 \quad \quad 2^2 \quad \quad 2^1 \quad \quad 2^0$$

Base 2

Most significant bit (MSB)

Least significant bit (LSB)

1 0 1 1

eights fours twos ones

$= 1 * 8 + 0 * 4 + 1 * 2 + 1 * 1 = 11_{10}$

Base 10 to Base 2

Question: What is 6 in base 2?

- Strategy:
  - What is the largest power of $2 \leq 6$? $\mathbf{2^2=4}$
  - Now, what is the largest power of $2 \leq 6 - 2^2$? $\mathbf{2^1=2}$
  - $6 - 2^2 - 2^1 = 0$

$$\underline{0} \quad \underline{1} \quad \underline{1} \quad \underline{0}$$
$$2^3 \quad\quad 2^2 \quad\quad 2^1 \quad\quad 2^0$$

$= \mathbf{0}*8 + \mathbf{1}*4 + \mathbf{1}*2 + \mathbf{0}*1 = 6$

Practice: Base 2 to Base 10

What is the base-2 value of 1010 in base-10?

a) 20

b) 101

c) 10

d) 5

e) Other

1010 isn't 1010 so much as it's 8 + 2

Practice: Base 10 to Base 2

What is the base-10 value of 14 in base 2?

a) 1111

b) 1110

c) 1010

d) Other

14 can be written as a sum of powers of 2

14 isn't 14 so much as it's 8 + 4 + 2

that can be encoded as 1110

Byte Values

What are the minimum and maximum base-10 values that a single byte (8 bits) can represent?

minimum = 0   **maximum = 255**

11111111

2^x: 7   6   5   4   3   2   1   0

- **Strategy 1:** $1 * 2^7 + 1 * 2^6 + 1 * 2^5 + 1 * 2^4 + 1 * 2^3 + 1 * 2^2 + 1 * 2^1 + 1 * 2^0 = 255$
- **Strategy 2:** $2^8 - 1 = 255$

Multiplying by Base

$$1450 \times 10 = 14500$$

$$1100_2 \times 10_2 = 11000$$

Key Idea: appending a 0 to the end effectively multiplies by the base!
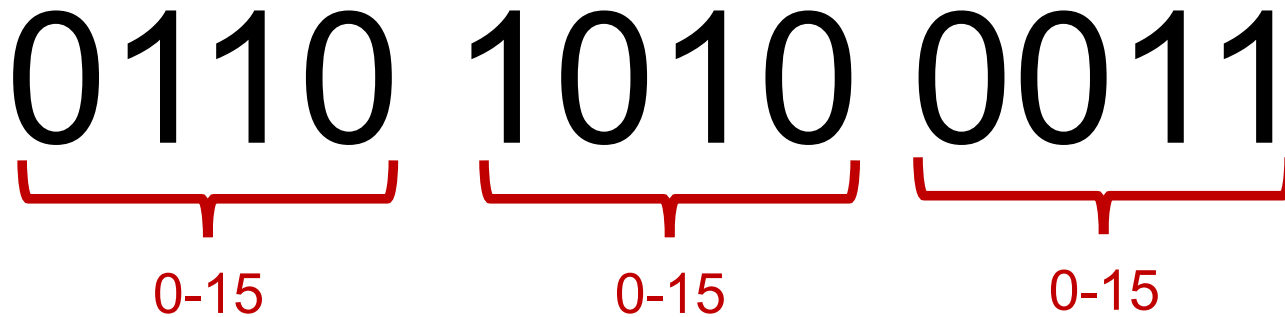
$$1450 / 10 = 145$$

$$1100_2 / 10_2 = 110$$

Key Idea: chomping off a 0 from the end divides by the base!

Hexadecimal

When working with 32- or 64-bit figures, binary representations get pretty long.

- Instead, we'll often encode numbers in **base-16**, or **hexadecimal**, instead.

$$0110 \quad 1010 \quad 0011$$

0-15   0-15   0-15

Hexadecimal

When working with 32- or 64-bit figures, binary representations get pretty long.

• Instead, we'll often encode numbers in **base-16**, or **hexadecimal**, instead**.**

0-15   0-15   0-15

Each quartet of bits can be rewritten as a single digit in **base-16**!

Hexadecimal

Hexadecimal is **base 16**, so we need digits for 0 through 15, inclusive. How?

0 1 2 3 4 5 6 7 8 9

Hexadecimal

- If it's not clear from context, we can explicitly identify numbers as hexadecimal by prefixing them with **0x** and identify numbers as binary using **0b**.

- e.g., **0xf5** is **0b11110101** is decimal number **245**

$$0x\underbrace{f}_{1111}\ \underbrace{5}_{0101}$$

Practice: Hexadecimal to Binary

What is **0x173A** in binary?

Hexadecimal	1	7	3	A
Binary	0001	0111	0011	1010

0x173A = 0b1011100111010

Practice: Binary to Hexadecimal

What is **0b1111001010** in hexadecimal? (*Hint: start from the right)*

Binary	11	1100	1010
Hexadecimal	3	C	A

0b1111001010 = 0x3CA

Hexadecimal: Quirky but concise

- Let's look at a single byte (8 bits):

165   base-10: Human-readable,
but cannot easily interpret on/off bits

0b10100101   base-2: Computers love this,
but most humans do not.

0xa5   base-16: Easy to convert to base-2,
More easily digested format
(fun fact: a half-byte is called a nibble.. tee hee hee)

Number Representations

- **Unsigned Integers**: positive integers and 0. (e.g., 0, 1, 2, … 99999…)

- **Signed Integers:** negative, positive integers and 0. (e.g., …-2, -1, 0, 1,… 9999…)

- **Floating Point Numbers:** real numbers. (e,g. 0.1, -12.2, 1.5×10^{12})

  → **Look up IEEE floating point if you're interested!**

Number Representations

C Declaration	Size (in bytes)
`int`	4
`double`	8
`float`	4
`char`	1
`char *`	8
`short`	2
`long`	8

Back When Jerry Learned C

C Declaration	Size (Bytes)
int	4
double	8
float	4
char	1
char *	4
short	2
long	4

Transitioning To Larger Data Types

- **Early 2000s:** most computers were **32-bit.** That meant pointers were **4 bytes**, or **32 bits.**

- 32-bit pointers store a memory address from 0 to 2^{32} - 1, equaling **2^{32} bytes of addressable memory**. That's **4 gigabytes**, meaning that 32-bit computers could address *at most* **4GB** of memory! That was a lot then, but it's not today.

- Most computers now are **64-bit.** Many data types got more memory, and pointers now need to be **64 bits.**

- 64-bit pointers can distinguish between addresses 0 to 2^{64} - 1, equaling **2^{64} bytes of addressable memory**. This equals **16 exabytes**, meaning that 64-bit computers could address up to **16 * 1024 * 1024 * 1024 GB** of memory!

Unsigned Integers

- An **unsigned** integer is either 0 or some positive integer (no negatives).

- We have already discussed the conversion between decimal and binary. Examples:

```
0b0001 = 1
0b0101 = 5
0b1011 = 11
0b1111 = 15
```

- The range of an unsigned number is set as $0 \rightarrow 2^w - 1$, where w is the number of bits—e.g., a 32-bit integer can represent 0 to $2^{32} - 1$, or 4,294,967,295.

Unsigned Integers

4-bit unsigned integer representation

Signed Integers

A **signed** integer is a negative integer, 0, or a positive integer.

- *Problem:* How can we represent negative *and* positive numbers in binary?

Idea: let the **most** significant bit represent sign and let the others represent magnitude.
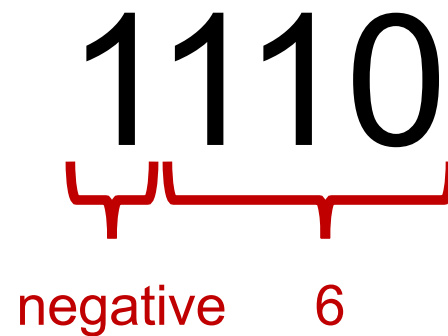
Sign Magnitude Representation: 4-bit

0110

positive   6

1110

negative   6

Sign Magnitude Representation: 4-bit

0011

positive   3

1011

negative   3

Sign Magnitude Representation: 4-bit

0000

positive   0

1000

negative   0

Sign Magnitude Representation: 4-bit

1 000 = -0	0 000 = 0
1 001 = -1	0 001 = 1
1 010 = -2	0 010 = 2
1 011 = -3	0 011 = 3
1 100 = -4	0 100 = 4
1 101 = -5	0 101 = 5
1 110 = -6	0 110 = 6
1 111 = -7	0 111 = 7

We're only representing 15 different values via 16 different patterns.
#sadness

Sign Magnitude Representation

- **Pro:** easy to represent, and easy to convert to and from decimal.
- **Con:** +/-0 is
- **Con:** we lose a bit that could be used to represent more numbers
- **Con:** arithmetic is tricky: we need to find the sign, perhaps subtract (borrow and carry, etc.), maybe change the sign, maybe not. This complicates how hardware implements something as fundamental as addition. This is the disadvantage we really care about.

Can we do better?

A Better Idea

Ideally, binary addition should work the same whether the numbers are positive or negative.

$$
\begin{array}{r}
0101 \\
+\,\color{red}{????} \\
\hline
0000
\end{array}
$$

A Better Idea

Ideally, binary addition should work the same whether the numbers are positive or negative.

$$0101$$
$$+1011$$
$$\overline{0000}$$

yes, it's really 10000, but we don't have a fifth bit to store that

leftmost one, so we let go of it

A Better Idea

Ideally, binary addition should work the same whether the numbers are positive or negative.

$$
\begin{array}{r}
0011 \\
+\ ???? \\
\hline
0000
\end{array}
$$

Ideally, binary addition should work whether the numbers are positive or negative.

$$\begin{array}{r} 0011 \\ +1101 \\ \hline 0000 \end{array}$$

Ideally, binary addition should work the same whether the numbers are positive or negative.

$$
\begin{array}{r}
0000 \\
+\ ???? \\
\hline
0000
\end{array}
$$

A Better Idea

Ideally, binary addition should work the same whether the numbers are positive or negative.

$$
\begin{array}{r}
0000 \\
+\ 0000 \\
\hline
0000
\end{array}
$$

There Seems To Be A Pattern

0101   0011   0000
+1011 +1101 +0000
───── ───── ─────
0000   0000   0000

The negated number is the original number **bitwise inverted, plus one more!**

There Seems To Be A Pattern

A binary number plus its inverse is all 1s.

$$0101$$
$$+1010$$
$$1111$$

Add 1 to this to carry over all 1s and get 0!

$$1111$$
$$+0001$$
$$0000$$

Two's Complement

- With **two's complement**, we represent a positive number as **itself**, and its negative counterpart as its **two's complement.**

- The **two's complement** of a number is the binary digits inverted, plus 1.

- This works to convert from positive to negative, **and** back from negative to positive!

4-bit two's complement signed integer representation

History: Two's complement

- Binary representation was first proposed by John von Neumann in *First Draft of a Report on the EDVAC* (1945).
  - That same year, he invented the merge sort algorithm

- Many early computers used either sign-magnitude or one's complement.

```
+7      0b0000 0111
-7      0b1111 1000
```
8-bit one's complement

- The System/360, developed by IBM in 1964, was widely popular—it had 1024KB memory!!!—and established two's complement as the dominant binary representation of integers.


EDVAC (1945)

System/360 (1964)

Two's Complement

- **Con:** more difficult to represent, and difficult to convert to and from decimal, between positive and negative.

- **Pro:** only 1 representation for 0. 😍

- **Pro:** the most significant bit still indicates the sign of a number.

- **Pro:** addition works for any combination of positive and negative, and electrical engineers love this.

4-bit two's complement signed integer representation

Two's Complement

Adding two numbers is just that: adding! And there is no special case for negative numbers. e.g., what is 2 + -5?

$$
\begin{array}{r}
0010 \\
+\,1011 \\
\hline
1101
\end{array}
$$

2

-5

-3

Two's Complement

Subtracting one number from a second is the same as adding the two's complement of that number from the second, e.g., 4 - 5 is just 4 + (-5).

$$
\begin{array}{r}
0100 \\
-\ 0101 \\
\hline
\end{array}
\qquad
\begin{array}{r}
4 \\
5 \\
\end{array}
\qquad\longrightarrow\qquad
\begin{array}{r}
0100 \\
+\ 1011 \\
\hline
1111 \\
\end{array}
\qquad
\begin{array}{r}
4 \\
-5 \\
-1 \\
\end{array}
$$