



# **CS107 Lecture 4**

## **Bits and Bytes Wrap-up, Bitwise Operators**

Reading: Bryant & O'Hallaron, Ch. 2.1

[Ed Discussion](#)

# Casting

What happens at the byte level when we cast between variable types? **The bytes remain the same! This means they may be interpreted differently depending on the type.**



```
int v = -12345;
unsigned int uv = v;
printf("v = %d, uv = %u\n", v, uv);
```

This prints out: "v = -12345, uv = 4294954951". **Why?**

The bit representation for -12345 is  
0b11111111111111111111111100111111000111.

If we treat this as an unsigned, inherently positive number, it's *huge*!

# Casting



# Casting

You can cast something to another type by putting that type in parentheses in front of the value:

```
int v = -12345;  
...(unsigned int)v...
```

You can also use the **U** suffix after a number literal to treat it as unsigned:

```
-12345U
```

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes



# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>	Signed	true	yes

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>	Signed	true	yes
<code>2147483647U &gt; -2147483648</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>	Signed	true	yes
<code>2147483647U &gt; -2147483648</code>	Unsigned	false	No!

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>	Signed	true	yes
<code>2147483647U &gt; -2147483648</code>	Unsigned	false	No!
<code>-1 &gt; -2</code>			
<code>(unsigned)-1 &gt; -2</code>			

# Comparisons Between Different Types

**Be careful** when comparing signed and unsigned integers. **C will implicitly cast** the signed argument to unsigned and then evaluate the expression assuming both numbers are unsigned and nonnegative.

Expression	Comparison Type?	Evaluates To?	Mathematically correct?
<code>0 == 0U</code>	Unsigned	true	yes
<code>-1 &lt; 0</code>	Signed	true	yes
<code>-1 &lt; 0U</code>	Unsigned	false	No!
<code>2147483647 &gt; -2147483648</code>	Signed	true	yes
<code>2147483647U &gt; -2147483648</code>	Unsigned	false	No!
<code>-1 &gt; -2</code>	Signed	true	yes
<code>(unsigned)-1 &gt; -2</code>	Unsigned	true	yes



# Expanding Bit Representations

- Sometimes, we need to convert between two integers of different sizes (e.g. **short** to **int**, or **int** to **long**).
- We might not be able to convert from a bigger data type to a smaller data type and retain all information, but we should always be able to convert from a **smaller** data type to a **larger** data type.
- For **unsigned** values, we can prepend *leading zeros* to the representation ("zero extension")
- For **signed** values, we can *repeat the sign of the value* for new digits ("sign extension")
- Note: when doing  $<$ ,  $>$ ,  $<=$ ,  $>=$  comparison between different size types, it will *promote the smaller type to the larger one*.

# Expanding Bit Representation

```
unsigned short s = 4;  
// short is a 16-bit format, so          s = 0000 0000 0000 0100b  
  
unsigned int i = s;  
// conversion to 32-bit int, so i = 0000 0000 0000 0000 0000 0000 0000 0100b
```

# Expanding Bit Representation

```
short s = 4;  
// short is a 16-bit format, so          s = 0000 0000 0000 0100b
```

```
int i = s;  
// conversion to 32-bit int, so i = 0000 0000 0000 0000 0000 0000 0000 0100b
```

— or —

```
short s = -4;  
// short is a 16-bit format, so          s = 1111 1111 1111 1100b
```

```
int i = s;  
// conversion to 32-bit int, so i = 1111 1111 1111 1111 1111 1111 1111 1100b
```

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
int x = 53191;  
short sx = x;  
int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit int), 53191:

**0000 0000 0000 0000 1100 1111 1100 0111**

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

**1100 1111 1100 0111**

This is -12345! And when we cast sx back an int, we sign-extend the number.

**1111 1111 1111 1111 1100 1111 1100 0111**     // still -12345

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
int x = -3;  
short sx = x;  
int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit int), -3:

**1111 1111 1111 1111 1111 1111 1111 1101**

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

**1111 1111 1111 1101**

This is -3! **If the number does fit, it will convert fine.** y looks like this:

**1111 1111 1111 1111 1111 1111 1111 1101** // still -3

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
unsigned int x = 128000;  
unsigned short sx = x;  
unsigned int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit unsigned int), 128000:

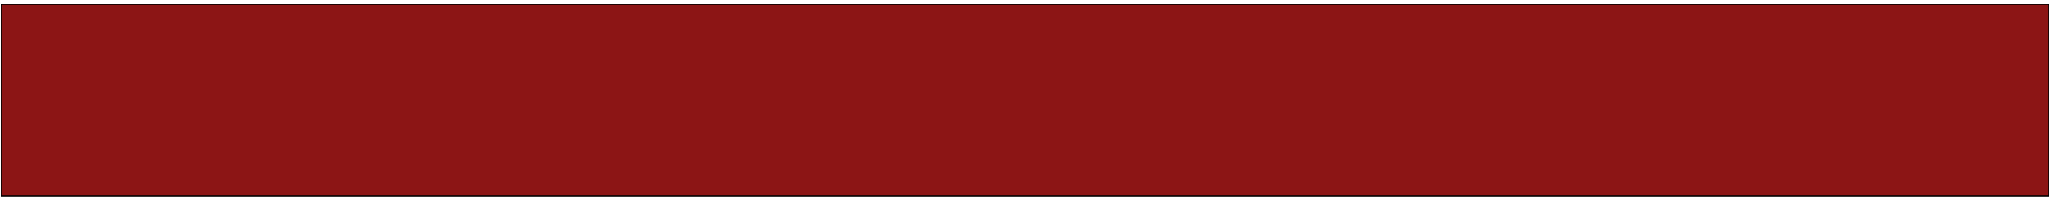
**0000 0000 0000 0001 1111 0100 0000 0000**

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

**1111 0100 0000 0000**

This is 62464! **Unsigned numbers can lose info too.** Here is what y looks like:

**0000 0000 0000 0000 1111 0100 0000 0000**     // still 62464



**Now that we understand  
values are really stored in  
binary, how can we  
manipulate them at the bit  
level?**

# Bitwise Operators

- You're already familiar with many operators in C:
  - **Arithmetic operators:** +, -, \*, /, %
  - **Comparison operators:** ==, !=, <, >, <=, >=
  - **Logical Operators:** &&, ||, !
- Today and Wednesday, we'll be discussing a new category of operators:  
**bitwise operators:**

&, |, ~, ^, <<, >>



# And (&)

AND is a binary operator. The AND of 2 bits is 1 if both bits are 1, and 0 otherwise.

**output = a & b;**

a	b	output
0	0	0
0	1	0
1	0	0
1	1	1

& with 1 to let a bit through, & with 0 to zero out a bit

# Or (|)

OR is a binary operator. The OR of 2 bits is 1 if either (or both) bits is 1.

**output = a | b;**

a	b	output
0	0	0
0	1	1
1	0	1
1	1	1

| with 1 to turn on a bit, | with 0 to let a bit go through

# Not ( $\sim$ )

NOT is a unary operator. The NOT of a bit is 1 if the bit is 0, or 0 otherwise.

**output =  $\sim$ a;**

a	output
0	1
1	0

# Exclusive Or (^)

Exclusive Or (XOR) is a binary operator. The XOR of 2 bits is 1 if *exactly* one of the bits is 1, or 0 otherwise.

**output = a ^ b;**

a	b	output
0	0	0
0	1	1
1	0	1
1	1	0

^ with 1 to flip a bit, ^ with 0 to let a bit go through unmodified

# Operators on Multiple Bits

- When these operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:

AND

	0110
&	1100
	----
	0100

OR

	0110
	1100
	----
	1110

XOR

	0110
^	1100
	----
	1010

NOT

~	1100
	----
	0011

# Operators on Multiple Bits

- When these operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:


AND	OR	XOR	NOT
<pre>0110 &amp; 1100 ---- 0100</pre>	<pre>0110   1100 ---- 1110</pre>	<pre>0110 ^ 1100 ---- 1010</pre>	<pre>~ 1100 ---- 0011</pre>

**Note:** these are different from the logical operators AND (&&), OR (||) and NOT (!).

# Operators on Multiple Bits

- When these operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:

AND	OR	XOR	NOT
<pre>0110 &amp; 1100 ---- 0100</pre>	<pre>0110   1100 ---- 1110</pre>	<pre>0110 ^ 1100 ---- 1010</pre>	<pre>~ 1100 ---- 0011</pre>



This is different from logical AND (&&). The logical AND returns true if both are nonzero, or false otherwise. With &&, this would be 6 && 12, which would evaluate to **true** (1).

# Operators on Multiple Bits

- When these operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:

AND	OR	XOR	NOT
<pre>0110 &amp; 1100 ---- 0100</pre>	<pre>0110   1100 ---- 1110</pre>	<pre>0110 ^ 1100 ---- 1010</pre>	<pre>~ 1100 ---- 0011</pre>

This is different from logical OR (`||`). The logical OR returns true if either are nonzero, or false otherwise. With `||`, this would be `6 || 12`, which would evaluate to **true** (1).



# Operators on Multiple Bits

- When these operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:

AND	OR	XOR	NOT
<pre>0110 &amp; 1100 ---- 0100</pre>	<pre>0110   1100 ---- 1110</pre>	<pre>0110 ^ 1100 ---- 1010</pre>	<pre>~ 1100 ---- 0011</pre>

This is different from logical NOT (!). The logical NOT returns true if this is zero, and false otherwise. With !, this would be !12, which would evaluate to **false** (0).

# Bitmasks

We will frequently want to manipulate or otherwise isolate specific bits in a larger collection of them. A **bitmask** is a constructed bit pattern that we can use, along with standard bit operators like `&`, `|`, `^`, `~`, `<<`, and `>>`, to do this.

**Motivating Example:** Bit vectors

**Aside:** C++ relies on bit vectors to efficiently implement **`vector<bool>`**.

# Bit Vectors and Sets

Instead of using arrays of Booleans, one can more compactly store Boolean information in bits instead.

- **Example:** we can represent current courses taken using a **char** and manipulate its contents using bit operators.

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

# Bit Vectors and Sets

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

- How do we find the union of two sets of courses taken? Use OR:

```
00100011
| 01100001
-----
01100011
```

# Bit Vectors and Sets

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

- How do we find the intersection of two sets of courses taken? Use AND:

```
00100011
& 01100001
-----
00100001
```

# Bit Masking

**Example:** how do we update our bit vector to indicate we've taken CS107?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

```
  00100011
| 00001000
-----
  00101011
```

# Bit Masking

```
#define CS106A 0x1    /* 0000 0001 */
#define CS106B 0x2    /* 0000 0010, or 0x1 << 1 */
#define CS106AX 0x4   /* 0000 0100, or 0x1 << 2 */
#define CS107 0x8     /* 0000 1000, or 0x1 << 3 */
#define CS111 0x10    /* 0001 0000, or 0x1 << 4 */
#define CS103 0x20    /* 0010 0000, or 0x1 << 5 */
#define CS109 0x40    /* 0100 0000, or 0x1 << 6 */
#define CS161 0x80    /* 1000 0000, or 0x1 << 7 */

char myClasses = ...;
myClasses = myClasses | CS107;    // include CS107!
```

# Bit Masking

```
#define CS106A 0x1    /* 0000 0001 */
#define CS106B 0x2    /* 0000 0010, or 0x1 << 1 */
#define CS106AX 0x4   /* 0000 0100, or 0x1 << 2 */
#define CS107 0x8     /* 0000 1000, or 0x1 << 3 */
#define CS111 0x10    /* 0001 0000, or 0x1 << 4 */
#define CS103 0x20    /* 0010 0000, or 0x1 << 5 */
#define CS109 0x40    /* 0100 0000, or 0x1 << 6 */
#define CS161 0x80    /* 1000 0000, or 0x1 << 7 */

char myClasses = ...;
myClasses |= CS107;    // include CS107!
```



# Bit Masking

- **Example:** how do we update our bit vector to indicate we've *dropped* CS103?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

```
00100011
& 11011111
-----
00000011
```

```
char myClasses = ...;
myClasses &= ~CS103;    // Drop CS103
```

# Bit Masking

- **Example:** how do we check if we've taken CS106B?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS111	CS107	CS106AX	CS106B	CS106A

```
00100011
& 00000010
-----
00000010
```

```
char myClasses = ...;
if (myClasses & CS106B) {...
    // taken CS106B!
```

# Bitwise Operator Tricks

- `|` with 1 is useful for turning select bits on
- `&` with 0 is useful for turning select bits off
- `|` is useful for taking the union of bits
- `&` is useful for taking the intersection of bits
- `^` is useful for flipping isolated bits
- `~` is useful for flipping all bits