CS107, Lecture 10 Stack and Heap

Reading: K&R 5.6-5.9 or Essential C section 6 on the heap Ed Discussion

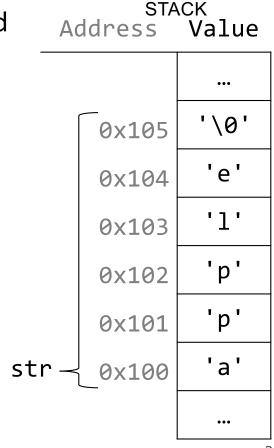
Arrays

When you declare an array, contiguous memory is allocated on the stack to store the contents of the entire array.

```
char str[6];
strcpy(str, "apple");
```

The array variable (e.g. **str**) is not a pointer; it refers to the entire array contents. In fact, **sizeof** returns the size of the entire array!

```
size_t num_bytes = sizeof(str); // 6
```



Arrays as Parameters

```
When you pass an array as a parameter, C makes a
                                                              Address
                                                                        Value
copy of the address of the first array element, and
                                                                  0x1f2
                                                                          '\0'
passes it (a pointer) to the function.
                                                                  0x1f1
                                                                          'h'
                                               main()
                                                          str.
                                                                  0x1f0
void myfunc(char *mystr) {
                                                                   0xff
                                                                   0xfe
                                                                   0xfd
int main(int argc, char *argv[]) {
                                                                   0xfc
      char str[3];
                                                                        0x1f0
                                            myfunc()
                                                                   0xfb
      strcpy(str, "hi");
                                                                   0xfa
      myfunc(str);
                                                                   0xf9
                                                             mystr 0xf8
```

Arrays as Parameters


```
This also means we can no longer get the full size of
                                                           Address
                                                                      Value
the array using sizeof, because now it is just a
                                                                       '\0'
                                                                0x1f2
pointer.
                                                                0x1f1
                                                                       'h'
                                             main()
                                                        str.
                                                                0x1f0
void myfunc(char *myStr) {
     size t size = sizeof(myStr); // 8
                                                                0xff
                                                                0xfe
                                                                0xfd
int main(int argc, char *argv[]) {
                                                                0xfc
     char str[3];
                                                                     0x1f0
                                           myfunc()
                                                                0xfb
     strcpy(str, "hi");
                                                                0xfa
      size t size = sizeof(str); // 3
     myfunc(str);
                                                                0xf9
                                                          mystr 0xf8
```

sizeof returns the size of an array, or 8 for a pointer. Therefore, when we pass an array as a parameter, we can no longer use sizeof to get its full size.


When you do pointer arithmetic, you are adjusting the pointer by a certain *number of places* (e.g., characters).

DATA SEGMENT Address Value '\0' 0xff5 ' ۾ ' 0xff4 '1' 0xff3 0xff2 0xff1 0xff0

Pointer arithmetic does *not* work in bytes. Instead, it works in the *size of the type it points to*.



Pointer arithmetic does *not* work in bytes. Instead, it works in the *size of the type it points to*.



When you use bracket notation with a pointer, you are actually *performing pointer arithmetic and dereferencing*:

```
char *str = "apple";  // e.g. 0xff0

// both of these add four places to str,

// and then dereference to get the char there.

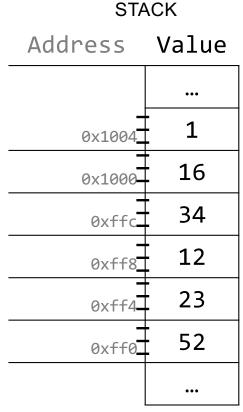
// e.g. get memory at 0xff4.

char fifth = str[4];  // 'e'

char fifth = *(str + 4);  // 'e'
```

DATA SEGMENT Address Value '\0' 0xff5 ' ۾ ' 0xff4 '1' 0xff3 0xff2 0xff1 0xff0

Pointer arithmetic with two pointers does *not* give the byte difference. Instead, it counts the number of **quantum elements** in between the two addresses.



CS107 Topic 3: How can we effectively manage all types of memory in our programs?

CS107 Topic 3

How can we effectively manage all types of memory in our programs?

Why is answering this question important?

- Shows us how we can pass around data efficiently using pointers (last time)
- Introduces us to the heap and the allocation of memory that we manage ourselves (this time)

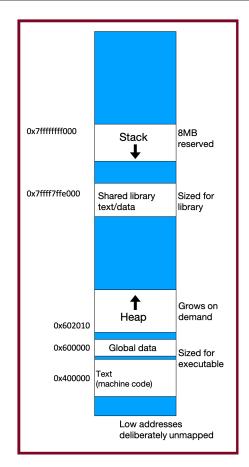
assign3: implement a function using resizable arrays to read lines of any length from a file and write 2 programs using that function to print the last N lines of a file and print just the unique lines of a file. These programs emulate the **tail** and **uniq** Unix builtins.

Learning Goals

- Learn about the differences between the stack and the heap and when to use each one.
- Become familiar with the **malloc**, **realloc** and **free** functions for managing memory on the heap.

Memory Layout

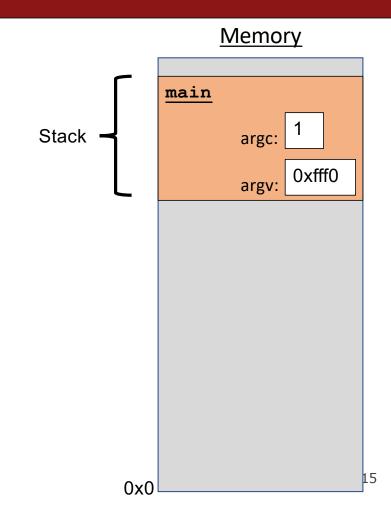
- We are going to dive deeper into different areas of memory used by our programs.
- The **stack** is the place where all local variables and parameters live for each function. A function's "stack frame" goes away when the function returns.
- The stack grows **downwards** when a new function is called and shrinks **upwards** when the function exits.



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

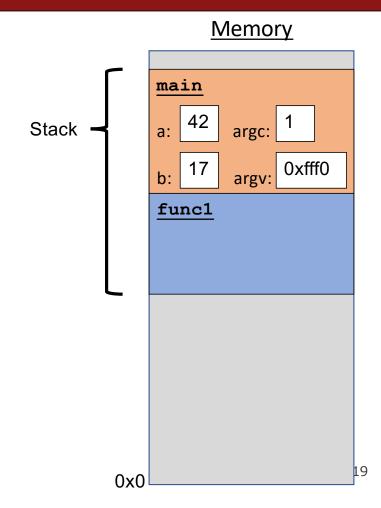
int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

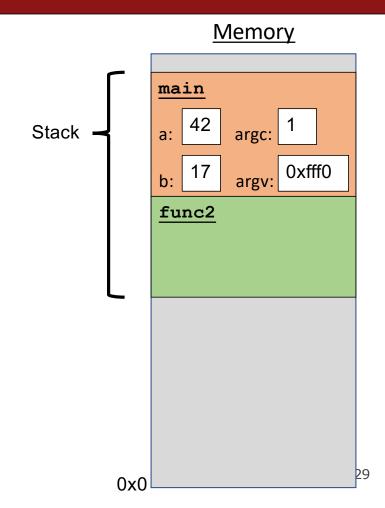
int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}


int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 argv: 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

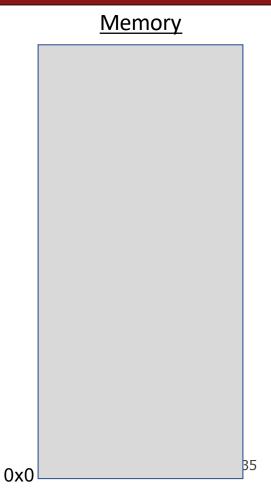
int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```

Memory main Stack argc: 0xfff0 0x0

```
void func2() {
    int d = 0;
}

void func1() {
    int c = 99;
    func2();
}

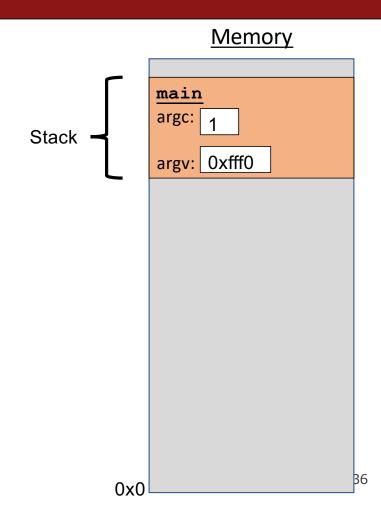
int main(int argc, char *argv[]) {
    int a = 42;
    int b = 17;
    func1();
    func2();
    printf("Done.");
    return 0;
}
```



The Stack Failing Us

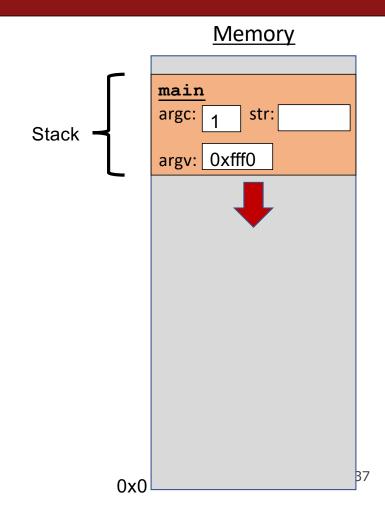
```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



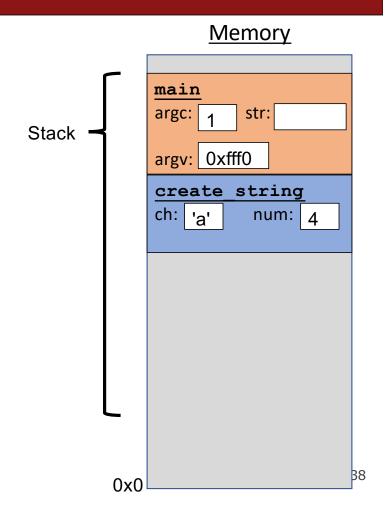
```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



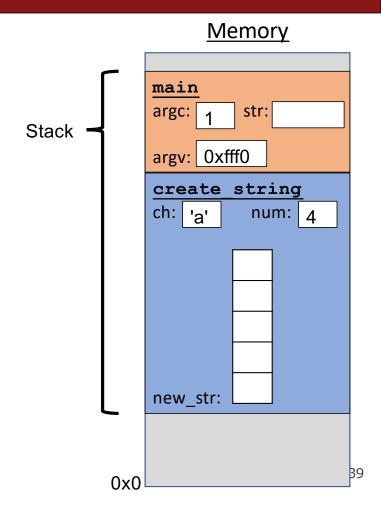
```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



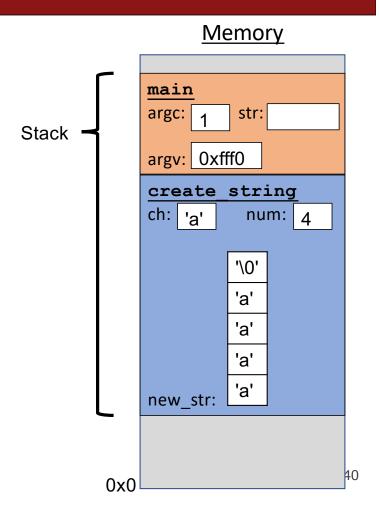
```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

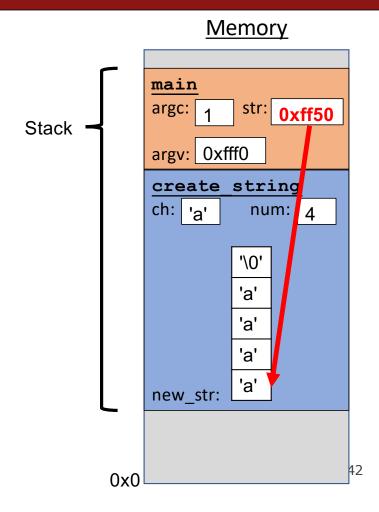
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
Memory
char *create_string(char ch, int num) {
    char new_str[num + 1];
                                                                   main
    for (int i = 0; i < num; i++) {
                                                                   argc:
                                                                           str:
        new str[i] = ch;
                                                        Stack
                                                                   argv: 0xfff0
    new_str[num] = '\0';
                                                                   create string
    return new str;
                                                                   ch: 'a'
                                                                            num: 4
                                                                           '\0'
                                            Returns e.g. 0xff50
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
                                                                           'a'
    printf("%s", str); // want "aaaa"
                                                                           'a'
    return 0;
                                                                           'a'
                                                                   new_str:
                                                                0x0
```

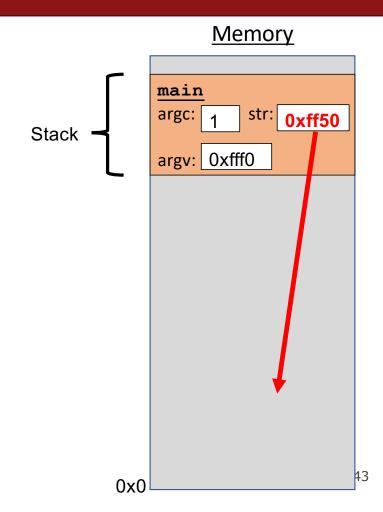
```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    new_str[num] = '\0';
    return new str;
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
    Ack! Local variables go away when the function declaring
    them exits. These 'a's were embedded in a variable that
     no longer exists, so str is not our address to deference.
```

Memory main 0xff50

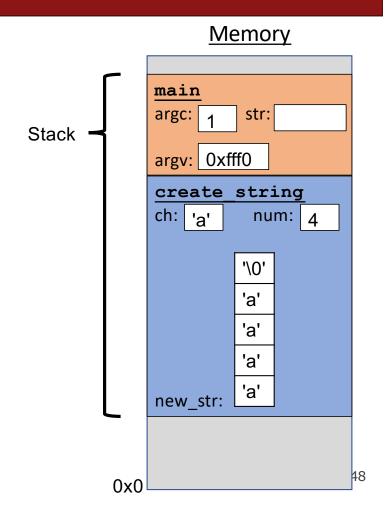
0x0

```
Memory
char *create_string(char ch, int num) {
    char new_str[num + 1];
                                                                main
    for (int i = 0; i < num; i++) {
                                                                          0xff50
        new_str[i] = ch;
    new_str[num] = '\0';
    return new str;
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
    Sometimes, we can make the array in the caller and pass
    it as a parameter. But this isn't always possible if the size
                    isn't known in advance.
                                                            0x0
```

This is a problem! We need a way to allocate memory that persists even after the allocating function exits.

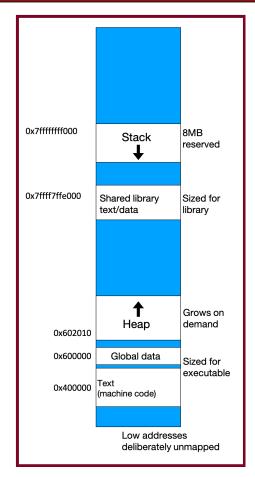
```
Memory
char *create_string(char ch, int num) {
    char new_str[num + 1];
                                                                    main
    for (int i = 0; i < num; i++) {
                                                                    argc:
                                                                            str:
         new str[i] = ch;
                                                        Stack
                                                                    argv: 0xfff0
    new_str[num] = '\0';
                                                                    create string
    return new str;
                                                                    ch: 'a'
                                                                            num: 4
                                                                            '\0'
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
                                                                            'a'
    printf("%s", str); // want "aaaa"
                                                                            'a'
    return 0;
                                                                            'a'
                   Us: Hey C, is there a way to allocate this
                                                                    new_str:
                   variable so it persists beyond the lifetime
                        of the function that allocates it?
                                                                0x0
```

```
char *create_string(char ch, int num) {
    char new_str[num + 1];
    for (int i = 0; i < num; i++) {
        new str[i] = ch;
    new_str[num] = '\0';
    return new str;
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
                    C: sure, but since I don't know
                   when to deallocate it, it's your
                       responsibility to do that
```



- The **heap** is a part of memory below the stack that you manage yourself. Unlike the stack, the memory only goes away when you deallocate it.
- Unlike the stack, the heap grows **upwards** as more memory is allocated.

The heap is **dynamic memory** – memory that can be allocated, resized, and freed during **program execution**.



Working with the heap

Working with the heap consists of 3 core steps:

- 1. Allocate memory with malloc/realloc/strdup/calloc
- 2. Assert heap pointer is not **NULL**
- 3. Free memory when done using **free**.

The heap provides **dynamic memory** that you programmatically introduce—sometimes incorrectly—to the program. That means you may encounter **runtime errors**, even if your code compiles! It's your responsibility to allocate properly and debug when there are problems.

malloc


```
void *malloc(size_t size);
```

To allocate memory on the heap, use the **malloc** function and specify the **number of bytes** you need.

- This function returns a pointer to the **leading address** of the new memory block. It doesn't know or care if it's to be used for an array, a struct, or anything else.
- void * denotes a pointer to generic memory. You can set another pointer equal to it without any casting.
- The memory is *not* zeroed out!
- If malloc returns NULL, the heap couldn't service the allocation request.

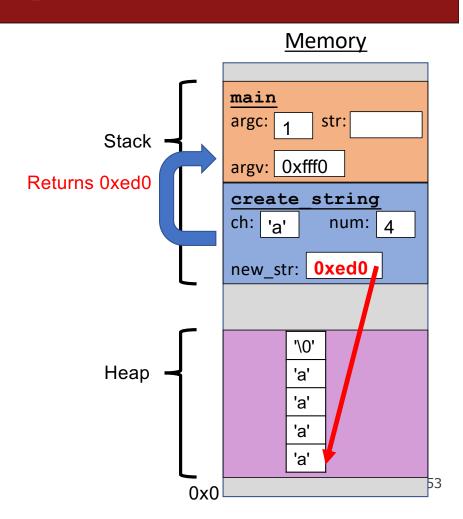
```
char *create_string(char ch, int num) {
    char *new_str = malloc(num + 1);
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



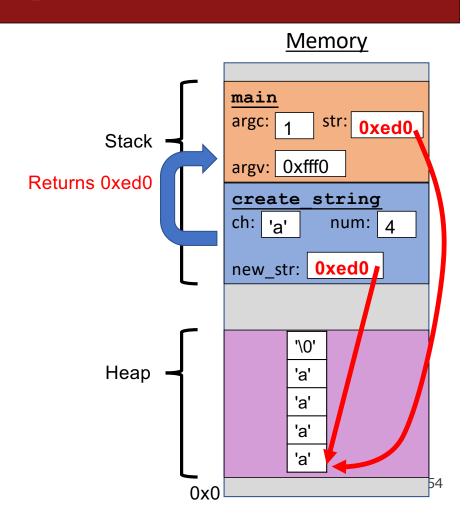
```
char *create_string(char ch, int num) {
    char *new_str = malloc(num + 1);
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



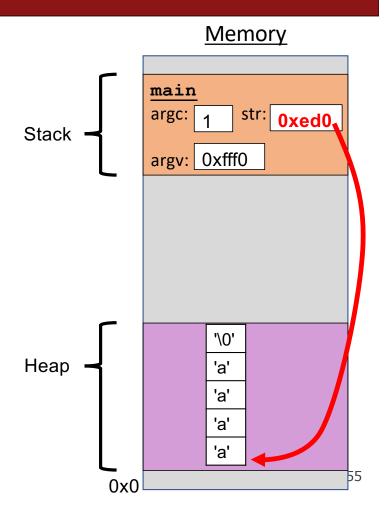
```
char *create_string(char ch, int num) {
    char *new_str = malloc(num + 1);
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



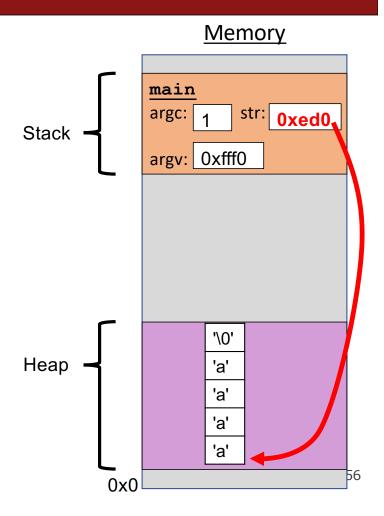
```
char *create_string(char ch, int num) {
    char *new_str = malloc(num + 1);
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
char *create_string(char ch, int num) {
    char *new_str = malloc(num + 1);
    for (int i = 0; i < num; i++) {
        new_str[i] = ch;
    }
    new_str[num] = '\0';
    return new_str;
}

int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
    return 0;
}</pre>
```



```
Memory
char *create_string(char ch, int num) {
    char *new_str = malloc(sizeof(char) * (num + 1));
                                                               main
    for (int i = 0; i < num; i++) {
                                                                      str: 0xed0
        new_str[i] = ch;
                                                    Stack
    new_str[num] = '\0';
    return new str;
int main(int argc, char *argv[]) {
    char *str = create_string('a', 4);
    printf("%s", str); // want "aaaa"
                                                                    '\0'
    return 0; // should free str, we will soon
                                                    Heap
                                                                    'a'
                                                                    'a'
                                                           0x0
```

Exercise: malloc multiples

Let's write a function that returns an array of the first len multiples of mult.

```
1 int *array_of_multiples(int mult, int len) {
2    /* TODO: arr declaration here */
3
4    for (int i = 0; i < len; i++) {
        arr[i] = mult * (i + 1);
    }
7    return arr;
8 }</pre>
```

Line 2: How should we declare arr?

```
A. int arr[len];
B. int arr[] = malloc(sizeof(int));
C. int *arr = malloc(sizeof(int) * len);
D. int *arr = malloc(sizeof(int) * (len + 1));
```



Exercise: malloc multiples

Let's write a function that returns an array of the first len multiples of mult.

```
int *array_of_multiples(int mult, int len) {
    /* TODO: arr declaration here */

for (int i = 0; i < len; i++) {
    arr[i] = mult * (i + 1);
    }
    return arr;
}</pre>
• Use a porter returned
• malloc's
```

Line 2: How should we declare arr?

```
A. int arr[len];
B. int arr[] = malloc(sizeof(int));
C. int *arr = malloc(sizeof(int) * len);
D. int *arr = malloc(sizeof(int) * (len + 1));
```

- Use a pointer to store the address returned by malloc.
- malloc's argument is the number of bytes to allocate.
- This code is missing an assertion.

Always assert with the heap

Let's write a function that returns an array of the first len multiples of mult.

```
int *array_of_multiples(int mult, int len) {
   int *arr = malloc(sizeof(int) * len);
   assert(arr != NULL);
   for (int i = 0; i < len; i++) {
      arr[i] = mult * (i + 1);
   }
   return arr;
}</pre>
```

- If an allocation error occurs (e.g., out of heap memory), malloc will return NULL. This is an important case to check **for robustness**.
- **assert** will intentionally end the program if the provided condition is false. A memory allocation error is significant, and we should terminate the program when we see them.

Other heap allocations: calloc

```
void *calloc(size_t nmemb, size_t size);
calloc is like malloc that zeros out the memory for you—thanks, calloc!
```

• You might notice its interface is also a little different—it takes two parameters, which are multiplied to calculate the number of bytes (nmemb * size).

```
// allocate and zero 20 ints
int *scores = calloc(20, sizeof(int));
// alternate (but slower)
int *scores = malloc(20 * sizeof(int));
for (int i = 0; i < 20; i++) scores[i] = 0;</pre>
```

• calloc is more expensive than malloc because it zeroes out all memory. Use only when absolutely necessary!

Other heap allocations: strdup

```
char *strdup(char *s);
```

strdup is a convenience function that returns a **null-terminated**, heapallocated string with the provided text, instead of requiring you to **malloc** and copy in the string yourself.

```
char *str = strdup("Hello, world!"); // on heap
str[0] = 'h';
```

You could imagine **strdup** might be implemented in terms of **malloc** + **strcpy**. (In fact, it pretty much is.)

Cleaning Up with free

```
void free(void *ptr);
```

- If we allocated memory on the heap and no longer need it, it is our responsibility to deallocate it.
- To do this, use the **free** command and pass in the *starting address on the* heap for the memory you no longer need.
- Example:

```
char *bytes = malloc(4);
...
free(bytes);
```

Free

```
void free(void *ptr);
```

When you free an allocation, you are freeing up what it *points* to. You are not deallocating the pointer itself. You can still use the pointer to point to something else.

```
char *str = strdup("hello");
...
free(str);
str = strdup("hi");
```



free details

Even if you have multiple pointers to the same block of memory, each memory block should only be freed **once**.

You must free the address you received in the previous allocation call. You cannot free just part of a previous allocation.

```
char *bytes = malloc(4);
char *ptr = malloc(10);
...
free(bytes);
...
free(ptr + 1);
```

Memory Leaks

A **memory leak** is when you do not free memory you previously allocated.

```
char *str = strdup("hello");
...
str = strdup("hi"); // memory leak! Lost previous str
```

Memory Leaks

- A memory leak occurs when you dynamically allocate a block of memory on the heap but fail to free it.
- Your program should be responsible for cleaning up any memory it allocates but no longer needs.
- If you never free any memory and allocate a large amount, you may run out of heap memory! (Running out of memory is rare, but it can happen if the program is designed to run for a very, long time—e.g., a web server.)
- However, memory leaks rarely cause crashes.
- We recommend not to worry about freeing memory until your program is written. Then, go back and free memory as appropriate.
- valgrind is a very helpful tool for finding memory leaks so they can be plugged.

realloc

```
void *realloc(void *ptr, size_t size);
```

- The **realloc** function takes an existing allocation pointer and enlarges to a new requested size. It returns the new pointer.
- If there is enough space after the existing memory block on the heap for the new size, **realloc** simply adds that space to the allocation.
- If there is not enough space, **realloc** moves the memory to a larger location, frees the old memory for you, and returns a pointer to the new location.

realloc

```
char *str = strdup("Hello");
assert(str != NULL);
...

// want to make str longer to hold "Hello world!"
char *addition = " world!";
str = realloc(str, strlen(str) + strlen(addition) + 1);
assert(str != NULL);

strcat(str, addition);
printf("%s", str);
free(str);
```

realloc

- realloc only accepts pointers that were previously returned by malloc/etc.
- Make sure to not pass pointers to the middle of heap-allocated memory.
- Make sure to not pass pointers to stack memory.

Cleaning Up with free and realloc

You only need to free the new memory coming out of realloc—the previous (smaller) one was already reclaimed by realloc.

```
char *str = strdup("Hello");
assert(str != NULL);
...
// want to make str longer to hold "Hello world!"
char *addition = " world!";
str = realloc(str, strlen(str) + strlen(addition) + 1);
assert(str != NULL);
strcat(str, addition);
printf("%s", str);
free(str);
```

Heap allocation interface: A summary

```
void *malloc(size_t size);
void *calloc(size_t nmemb, size_t size);
void *realloc(void *ptr, size_t size);
char *strdup(char *s);
void free(void *ptr);
```

Heap **memory allocation** guarantee:

- NULL on failure, so check with assert
- Memory is contiguous; it is not recycled unless you call free
- realloc preserves existing data
- calloc zero-initializes bytes, malloc and realloc do not

Undefined behavior occurs:

- If you overflow (i.e., you access beyond bytes allocated)
- If you use after free, or if free is called twice on a location.
- If you realloc/free non-heap address

Engineering principles: stack vs heap

Stack (for local variables)

Heap (dynamic memory)

- Fast
 Fast to allocate/deallocate; okay to oversize
- Convenient.
 Automatic allocation/ deallocation;
 declare/initialize in one step
- Reasonable type safety
 Thanks to the compiler
- Not especially plentiful Total stack size fixed, default 8MB
- Somewhat inflexible
 Cannot add/resize at runtime, scope
 dictated by control flow in/out of functions

Engineering principles: stack vs heap

Stack (for local variables)

- Fast
 Fast to allocate/deallocate; okay to oversize
- Convenient.
 Automatic allocation/ deallocation;
 declare/initialize in one step
- Reasonable type safety
 Thanks to the compiler
- Not especially plentiful

 Total stack size fixed, default 8MB
- Somewhat inflexible
 Cannot add/resize at runtime, scope
 dictated by control flow in/out of functions

Heap (dynamic memory)

- Plentiful.
 Can provide more memory on demand!
- Very flexible.
 Runtime decisions about how much/when to allocate, can resize easily with realloc
- Scope under programmer control
 Can precisely determine lifetime
- Lots of opportunity for error
 Low type safety, forget to allocate/free
 before done, allocate wrong size, etc.,
 Memory leaks (much less critical)

Stack and Heap

- Generally, unless a situation requires dynamic allocation, stack allocation is preferred. Often both techniques are used together in a program.
- Heap allocation is a necessity when:
 - you have a very large allocation that could blow out the stack
 - you need to control the memory's lifetime and/or memory must persist beyond a function call