



CS107, Lecture 17

Assembly: Arithmetic and Logic Wrap, Control Flow

Reading: B&O 3.5-3.6

[Ed Discussion](#)

Large Multiplication

- Multiplying 64-bit numbers can produce a 128-bit result. How does x86-64 support this with only 64-bit registers?
- If you specify two operands to **imul**, it multiplies them together and truncates it to fit in the second of the two 64-bit register operands.

`imul S, D` $D \leftarrow D * S$

- If you specify one operand, it multiplies that by **%rax**, and splits the product across **2** registers. It puts the high-order 64 bits in **%rdx** and the low-order 64 bits in **%rax**.

| Instruction | Effect | Description |
|----------------------|--|------------------------|
| <code>imulq S</code> | $R[\%rdx]:R[\%rax] \leftarrow S \times R[\%rax]$ | Signed full multiply |
| <code>mulq S</code> | $R[\%rdx]:R[\%rax] \leftarrow S \times R[\%rax]$ | Unsigned full multiply |

Division and Remainder

| Instruction | Effect | Description |
|----------------------|--|-----------------|
| <code>idivq S</code> | $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S;$ $R[\%rax] \leftarrow R[\%rdx]:R[\%rax] \div S$ | Signed divide |
| <code>divq S</code> | $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S;$ $R[\%rax] \leftarrow R[\%rdx]:R[\%rax] \div S$ | Unsigned divide |

- Terminology: **dividend / divisor = quotient with remainder**
- **x86-64** supports dividing up to a 128-bit value by a 64-bit value.
- The high-order 64 bits of the dividend need to be prepared and stored in **%rdx**, the low-order 64 bits in **%rax**. The divisor is the only listed operand.
- The quotient is stored in **%rax**, and the remainder in **%rdx**.

Division and Remainder

| Instruction | Effect | Description |
|----------------------|--|---------------------|
| <code>idivq S</code> | $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S;$ $R[\%rax] \leftarrow R[\%rdx]:R[\%rax] \div S$ | Signed divide |
| <code>divq S</code> | $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S;$ $R[\%rax] \leftarrow R[\%rdx]:R[\%rax] \div S$ | Unsigned divide |
| <code>cqto</code> | $R[\%rdx]:R[\%rax] \leftarrow \text{SignExtend}(R[\%rax])$ | Convert to oct word |

- Terminology: **dividend** / **divisor** = **quotient with remainder**
- The high-order 64 bits of the dividend need to be prepared and stored in `%rdx`, the low-order 64 bits in `%rax`. The divisor is the only listed operand.
- Most division uses only 64-bit dividends. The **`cqto`** instruction sign-extends the 64-bit value in `%rax` into `%rdx` to fill both registers with the dividend, as the division instruction expects.



Compiler Explorer Demo

<https://godbolt.org/z/4cT75M4nd>

Code Reference: full_divide

```
// Returns x/y, stores remainder in location stored in remainder_ptr
long full_divide(long x, long y, long *remainder_ptr) {
    long quotient = x / y;
    long remainder = x % y;
    *remainder_ptr = remainder;
    return quotient;
}
```

```
full_divide:
    movq %rdi, %rax
    movq %rdx, %rcx
    cqto
    idivq %rsi
    movq %rdx, (%rcx)
    ret
```

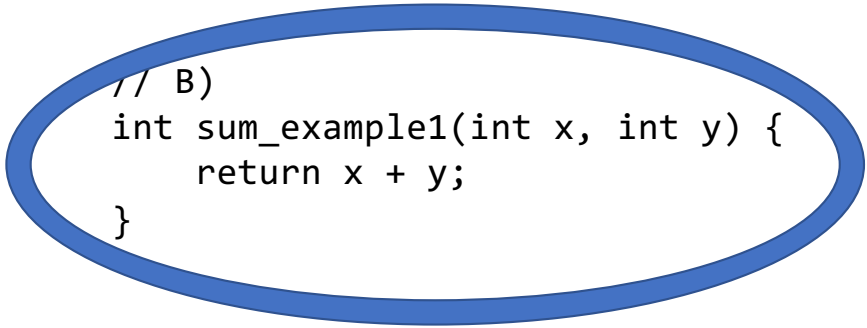
Assembly Exercise 1

```
000000000040116e <sum_example1>:  
40116e: 8d 04 37          lea    (%rdi,%rsi,1),%eax  
401171: c3                retq
```

Which of the following is most likely to have generated the above assembly?

```
// A)  
void sum_example1() {  
    int x;  
    int y;  
    int sum = x + y;  
}
```

```
// C)  
void sum_example1(int x, int y) {  
    int sum = x + y;  
}
```



```
// B)  
int sum_example1(int x, int y) {  
    return x + y;  
}
```

Assembly Exercise 2

```
0000000000401172 <sum_example2>:
    401172: 8b 47 0c          mov    0xc(%rdi),%eax
    401175: 03 07            add    (%rdi),%eax
    401177: 2b 47 18          sub    0x18(%rdi),%eax
    40117a: c3              retq
```

```
int sum_example2(int arr[]) {
    int sum = 0;
    sum += arr[0];
    sum += arr[3];
    sum -= arr[6];
    return sum;
}
```

What location or value in the assembly above represents the C code's **sum** variable?

%eax

Assembly Exercise 3

```
0000000000401172 <sum_example2>:
    401172: 8b 47 0c          mov    0xc(%rdi),%eax
    401175: 03 07            add    (%rdi),%eax
    401177: 2b 47 18          sub    0x18(%rdi),%eax
    40117a: c3               retq
```

```
int sum_example2(int arr[]) {
    int sum = 0;
    sum += arr[0];
    sum += arr[3];
    sum -= arr[6];
    return sum;
}
```

What location or value in the assembly code above represents the C code's **6** (as in **arr[6]**)?

0x18

Reverse Engineering 1

```
int add_to(int x, int arr[], int i) {  
    int sum = ____?____;  
    sum += arr[____?____];  
    return ____?____;  
}
```

```
// x in %edi, arr in %rsi, i in %edx  
add_to:  
    movslq %edx, %rdx  
    movl %edi, %eax  
    addl (%rsi,%rdx,4), %eax  
    ret
```

Reverse Engineering 1

```
int add_to(int x, int arr[], int i) {  
    int sum = ____?____;  
    sum += arr[____?____];  
    return ____?____;  
}
```

```
// x in %edi, arr in %rsi, i in %edx
```

```
add_to:
```

```
    movslq %edx, %rdx    // sign-extend i into full register  
    movl %edi, %eax      // copy x into %eax  
    addl (%rsi,%rdx,4), %eax // add arr[i] to %eax  
    ret
```

Reverse Engineering 1

```
int add_to(int x, int arr[], int i) {  
    int sum = x;  
    sum += arr[i];  
    return sum;  
}
```

```
// x in %edi, arr in %rsi, i in %edx
```

```
add_to:
```

```
    movslq %edx, %rdx          // sign-extend i into full register  
    movl %edi, %eax            // copy x into %eax  
    addl (%rsi,%rdx,4), %eax    // add arr[i] to %eax  
    ret
```

Reverse Engineering 2

```
int elem_arithmetic(int nums[], int y) {  
    int z = nums[___?___] * ___?___;  
    z -= ___?___;  
    z >>= ___?___;  
    return ___?___;  
}
```

```
// nums in %rdi, y in %esi
```

```
elem_arithmetic:
```

```
    movl %esi, %eax
```

```
    imull (%rdi), %eax
```

```
    subl 4(%rdi), %eax
```

```
    sarl $2, %eax
```

```
    addl $2, %eax
```

```
    ret
```

Reverse Engineering 2

```
int elem_arithmetic(int nums[], int y) {  
    int z = nums[___?___] * ___?___;  
    z -= ___?___;  
    z >>= ___?___;  
    return ___?___;  
}
```

```
// nums in %rdi, y in %esi
```

```
elem_arithmetic:
```

| | |
|---------------------------------|--|
| <code>movl %esi, %eax</code> | <code>// copy y into %eax</code> |
| <code>imull (%rdi), %eax</code> | <code>// multiply %eax by nums[0]</code> |
| <code>subl 4(%rdi), %eax</code> | <code>// subtract nums[1] from %eax</code> |
| <code>sarl \$2, %eax</code> | <code>// shift %eax right by 2</code> |
| <code>addl \$2, %eax</code> | <code>// add 2 to %eax</code> |
| <code>ret</code> | |

Reverse Engineering 2

```
int elem_arithmetic(int nums[], int y) {  
    int z = nums[0] * y;  
    z -= nums[1];  
    z >>= 2;  
    return z + 2;  
}
```

```
// nums in %rdi, y in %esi
```

```
elem_arithmetic:
```

| | |
|--------------------|-------------------------------|
| movl %esi, %eax | // copy y into %eax |
| imull (%rdi), %eax | // multiply %eax by nums[0] |
| subl 4(%rdi), %eax | // subtract nums[1] from %eax |
| sarl \$2, %eax | // shift %eax right by 2 |
| addl \$2, %eax | // add 2 to %eax |
| ret | |

Our First Assembly

```
int sum_array(int arr[], int nelems) {  
    int sum = 0;  
    for (int i = 0; i < nelems; i++) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

We're 1/2 of the way to understanding assembly!
What looks understandable right now?

0000000000401136 <sum_array>:

```
401136:  b8 00 00 00 00  
40113b:  ba 00 00 00 00  
401140:  39 f0  
401142:  7d 0b  
401144:  48 63 c8  
401147:  03 14 8f  
40114a:  83 c0 01  
40114d:  eb f1  
40114f:  89 d0  
401151:  c3
```

```
mov    $0x0,%eax  
mov    $0x0,%edx  
cmp    %esi,%eax  
jge    40114f <sum_array+0x19>  
movslq %eax,%rcx  
add    (%rdi,%rcx,4),%edx  
add    $0x1,%eax  
jmp    401140 <sum_array+0xa>  
mov    %edx,%eax  
retq
```



Executing Instructions

What does it mean for a program to execute?

Executing Instructions

So far:

- Program values can be stored in memory or registers.
- Assembly instructions read/write values back and forth between registers and main memory.
- Assembly instructions are **also stored in memory**.

Today:

- **Who controls the instructions?**
How do we know what to do now or next?

Answer:

- The **program counter**, stored in %rip.

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |



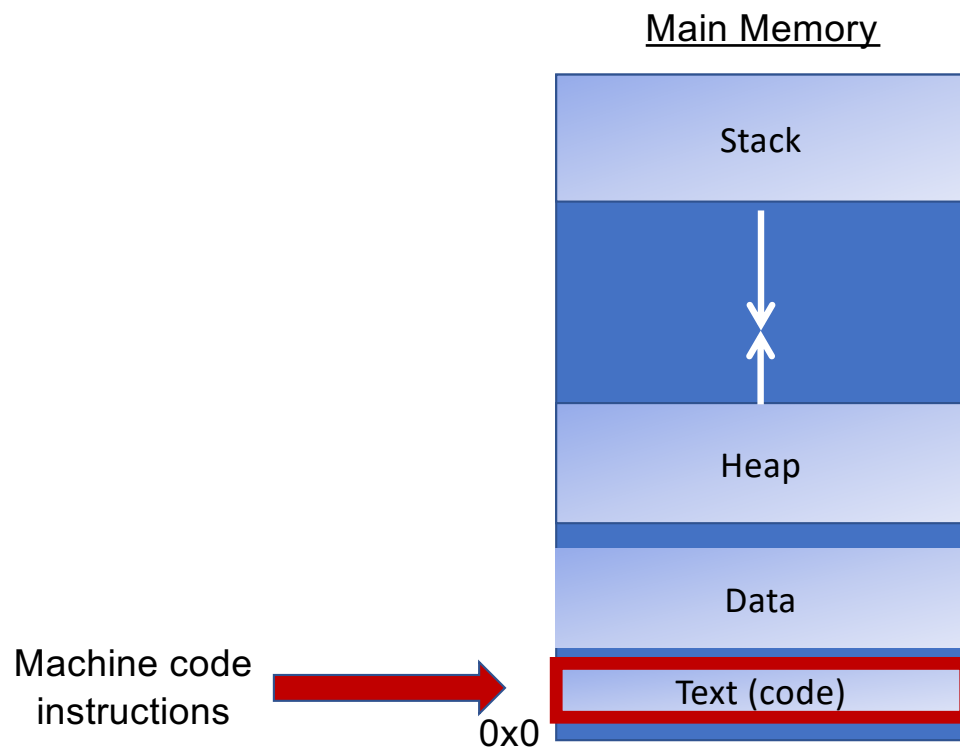
Register Responsibilities

Some registers take on special responsibilities during program execution.

- **%rax** stores the return value
- **%rdi** stores the first parameter to a function
- **%rsi** stores the second parameter to a function
- **%rdx** stores the third parameter to a function
- **%rip** stores the address of the next instruction to execute
- **%rsp** stores the address of the current top of the stack

See the x86-64 Guide and Reference Sheet on the Resources webpage for more!

Instructions Are Just Bytes!




%orig

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

00000000004004ed <loop>:

| | | |
|------------------------------|------|-------------------|
| 4004ed: 55 | push | %rbp |
| 4004ee: 48 89 e5 | mov | %rsp,%rbp |
| 4004f1: c7 45 fc 00 00 00 00 | movl | \$0x0,-0x4(%rbp) |
| 4004f8: 83 45 fc 01 | addl | \$0x1,-0x4(%rbp) |
| 4004fc: eb fa | jmp | 4004f8 <loop+0xb> |

Main Memory



%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

The **program counter** (PC), known as %rip in x86-64, stores the address in memory of the **next instruction** to be executed.

0x4004ed

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

The **program counter** (PC), known as %rip in x86-64, stores the address in memory of the *next instruction* to be executed.

0x4004ee

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

The **program counter** (PC), known as %rip in x86-64, stores the address in memory of the **next instruction** to be executed.

0x4004f1

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

The **program counter** (PC), known as %rip in x86-64, stores the address in memory of the **next instruction** to be executed.

0x4004f8

%rip

%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

The **program counter** (PC), known as %rip in x86-64, stores the address in memory of the **next instruction** to be executed.

0x4004fc

%rip

%rip

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00 movl \$0x0, -0x4(%rbp)

4004f8: 83 45 fc 01 addl \$0x1, -0x4(%rbp)

4004fc: eb fa jmp 4004f8 <loop+0xb>

Special hardware sets the program counter
to the next instruction:

%rip += size of bytes of current instruction

0x4004fc

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

Going In Circles

- How can we use this representation of execution to represent e.g., a **loop**?
- **Key Idea:** we can override what **%rip** stores and populate it with the address of an earlier instruction.

Jump!

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

0x4004fc

%rip

The **jmp** instruction is an **unconditional jump** that sets the program counter to the **jump target** (the operand).

Jump!

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

The **jmp** instruction is an **unconditional jump** that sets the program counter to the **jump target** (the operand).

0x4004fc

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

Jump!

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

0x4004fc

%rip

The **jmp** instruction is an **unconditional jump** that sets the program counter to the **jump target** (the operand).

Jump!

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

The **jmp** instruction is an **unconditional jump** that sets the program counter to the **jump target** (the operand).

0x4004fc

%rip

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

Jump!

00000000004004ed <loop>:

4004ed: 55

4004ee: 48 89 e5

4004f1: c7 45 fc 00 00 00 00

4004f8: 83 45 fc 01

4004fc: eb fa

push %rbp

mov %rsp,%rbp

movl \$0x0,-0x4(%rbp)

addl \$0x1,-0x4(%rbp)

jmp 4004f8 <loop+0xb>

| | |
|--------|----|
| 4004fd | fa |
| 4004fc | eb |
| 4004fb | 01 |
| 4004fa | fc |
| 4004f9 | 45 |
| 4004f8 | 83 |
| 4004f7 | 00 |
| 4004f6 | 00 |
| 4004f5 | 00 |
| 4004f4 | 00 |
| 4004f3 | fc |
| 4004f2 | 45 |
| 4004f1 | c7 |
| 4004f0 | e5 |
| 4004ef | 89 |
| 4004ee | 48 |
| 4004ed | 55 |

This assembly represents an infinite loop in C!

while (true) {...}

0x4004fc

%rip

jmp

The **jmp** instruction jumps to another instruction in the assembly code (an "unconditional jump").

| | |
|---------------------|------------------------|
| jmp Label | (Direct Jump) |
| jmp *Operand | (Indirect Jump) |

The destination can be hardcoded into the instruction (direct jump):

```
jmp 404f8 <loop+0xb> # jump to instruction at 0x404f8
```

The destination can also be one of the usual operand forms (indirect jump):

```
jmp *%rax # jump to instruction at address in %rax
```

"Interfering" with %rip

1. How do we repeat instructions in a loop?

`jmp [target]`

- A 1-step unconditional jump (always jump when we execute this instruction)

What if we want a **conditional jump**?